

线性方程组

本资料是微信公众号<机器人学家>编者的个人笔记，综合了各种线性代数课程、资料和自己的思考总结，仅供<机器人学家>公众号读者内部交流学习使用。

个人笔记难免有疏漏之处，关于笔记内容的任何疑问，可以在微信公众号<机器人学家>上与我直接交流，或者发送邮件给gorobotics@126.com，我尽量一一回复。



扫码关注微信公众号
机器人学家



扫码向我打赏
金额随意

目录

线性方程组

I. 解的存在性、唯一性

1.1 $AX=0$

1.2 $AX=b$

II. 矩阵分解factorization

2.1 LDU分解

2.1.1 Notes

2.1.2 Gaussian Elimination

Basic Idea

Record the process

2.1.3 解线性方程组

2.2 对角化

2.3 若当标准型

2.4 QR分解

2.4.1 One Algorithm

2.4.2 Notes on QR algorithms

2.5 SVD (重要)

2.5.1 idea

2.5.2 Summary

Matrix

Geometry

2.5.3 Solving linear systems

I. 解的存在性、唯一性

结合线性变换和线性空间理解。

1.1 $AX=0$

解即为A的null space。

有非零解 意味着A的null space不为空。

解性质：

- X的基础解系即Null (A) 的一组基。包括n-r组解。
- $AX = 0 \rightarrow BAX = 0$ ，所以BA对应的解包括A对应的解

求基础解系：把自由变量取成自然基

1.2 $AX=b$

X是A的列向量表出b的系数。

所以有解等价于b在A的row space里。

有唯一解等价于说A的null space is empty

解性质：齐次解+非齐次特解

这些都要结合空间理解。

II. 矩阵分解factorization

矩阵分解是为了让求解线性方程组更方便。

计算机中实际实现的线性方程组求解方法都是基于矩阵分解。

Some decomposition came from this idea(LDU, Jordan), while some others (SVD) have clear geometry meaning, when view A as a linear transformation.

2.1 LDU分解

A is $m \times n$, $Rank(A) = n$, $m \geq n$.

We could write $PA = LDU$, where:

P is a mxm permutation matrix. for example,

$$\begin{bmatrix} 0 & 0 & 1 \\ 1 & 0 & 0 \\ 0 & 1 & 0 \end{bmatrix}$$

L is a mxm square low-triangle matrix with 1s on the diagonal:

$$\begin{bmatrix} 1 & 0 & 0 \\ * & 1 & 0 \\ * & * & 1 \end{bmatrix}$$

D is a mxm square diagonal matrix.

U is a mxn upper-triangle matrix with 1s on the diagonal:

$$\begin{bmatrix} 1 & * & * \\ 0 & 1 & * \\ 0 & 0 & 1 \end{bmatrix}$$

We call the diagonal elements of D “pivot”

examples:

$$\begin{bmatrix} 1 & 0 \\ 1 & 1 \\ 0 & -1 \end{bmatrix} = \begin{bmatrix} 1 & 0 & 0 \\ 1 & 1 & 0 \\ 0 & -1 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 0 & 1 \\ 0 & 0 \end{bmatrix}$$

$$\begin{bmatrix} 1 & 1 & 0 \\ 2 & 1 & -1 \end{bmatrix} = \begin{bmatrix} 1 & 0 \\ 2 & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 \\ 0 & -1 \end{bmatrix} \begin{bmatrix} 1 & 1 & 0 \\ 0 & 1 & 1 \end{bmatrix}$$

2.1.1 Notes

P

The permutation matrix is used to change the ordering of base. We need to do this in case that a pivot is zero. PA changes row.

If needs to change coloum as well, we could write $P_1AP_2 = LDU$

U

Sometimes we treat D and U together as U, that is LU factorization: $PA = LU$

$\det(A) = \pm \det(D)$, sign depends on P.

if $A = A^T$ and $P = I$, then $U = L^T$

if $A = A^T$ and A is positive definite, Then we have:

$P = I, L = U^T, \det(A) = \det(D)$, diagonal entries of D are positive.

The LDU decomposition is **not unique**.

It is unique if:

1. A is square and invertable
2. $P=I$

2.1.2 Gaussian Elimination

Basic Idea

Perform row operations on A to obtain an Upper-triangular matrix.

Record the process

Two ways.

First one

While doing row operations on A, do the inverse, **row** operations on I

(Note: 初等矩阵 $E_{ij}(a)$ 左乘是 i 行的 a 倍加到 j 行, 右乘是 j 列的 a 倍加到 i 列)

To see what happend: $L_3 L_2 L_1 A = U$, then $A = L_1^{-1} L_2^{-1} L_3^{-1} U$

when $A \rightarrow U'$ (upper triangular), $I \rightarrow L$, L should be low-triangular.

then we have $A = LU'$.

When permutation is needed, say $L'A' = A$, A' needs to become $P_1 A'$.

$LP_1(P_1^{-1}A') = A$ the problem is that how could ensure LP_1 is still a lower-triangular matrix.

The solution is $(P_1 LP_1)(P_1^{-1}A') = P_1 A$ it could be verified that when doing Gaussian Elimination, $(P_1 LP_1)$ is still a lower-triangular matrix.

Second one

Doing row operations on $[I A]$, obtain $[M U']$, then we have $MA = U', A = M^{-1} U'$

This is already LU decomposition. Next, simply get DU from U' :

$$\begin{bmatrix} 1 & -2 & 1 \\ 0 & 4 & 1 \\ 0 & 0 & \frac{1}{4} \end{bmatrix} = \begin{bmatrix} 1 & \boxed{} & \boxed{} \\ \boxed{} & 4 & \boxed{} \\ \boxed{} & \boxed{} & \frac{1}{4} \end{bmatrix} \begin{bmatrix} 1 & -2 & 1 \\ \boxed{} & 1 & \frac{1}{4} \\ \boxed{} & \boxed{} & 1 \end{bmatrix}$$

2.1.3 解线性方程组

$AX = B \Rightarrow LDUX = B$, then we could obtain X by solving two simpler linear systems:

$Ly = B$, then

$DUx = y$

2.2 对角化

见笔记“矩阵运算”

2.3 若当标准型

The following two methods are based on $A^T A$

2.4 QR分解

Suppose A is a $m \times n$ matrix with independent columns, we can factor A as:

$$A = QR,$$

where $Q \in m \times n, R \in n \times n$

The columns of Q forms the orthonormal basis of A 's column space.

Which means $Q^T Q = I$.

(Note QQ^T may not be I , since A may not be square)

R is invertible and upper-triangular.

2.4.1 One Algorithm

First, form $A^T A$. This is a positive definite symmetric matrix.

(it is positive definite because the col of A are independent.)

Second, compute the LDU factorization of $A^T A$. you will get $A^T A = LDL^T$.

Finally, let $R = D^{1/2} L^T$, and $Q = AR^{-1}$

2.4.2 Notes on QR algorithms

Unfortunately, both of these straightforward algorithms have poor numerical stability. In practice one use a different algorithm.

2.5 SVD (重要)

SVD is a generalization of diagonalization.

2.5.1 idea

Diagonalization could not work nicely when:

1. you could not find n eigenvector as a basis for both input and output space (for example, A is not square, or Algebraic multiplicity $>$ geometric multiplicity for some eigen value)
2. some eigen values are complex.

In geometry, this means we could not choose a basis under which a dilation transformation could be found to be equivalence to A .

In this case, we would need to change basis for both input space and outspace, so as to form a dilation transformation.

The transformation matrix using B_1 for input space, B_2 for output space is denoted as ${}^{B_1}f_{B_2}$.

By definition, ${}^{B_1}f_{B_2} = A$ when $B_1 = B_2 = E$.

For example,

2.5.2 Summary

Matrix

Preparation

For Any $m \times n$ matrix A , we could form the following two decomposition:

$$A^T A = V D V^T$$

$$A A^T = U D' U^T$$

U is an $m \times m$ orthogonal matrix whose columns are the eigenvectors of $A A^T$.

The columns of U forms an orthogonal basis for **the whole output space**.

The **first k columns** of U forms an orthogonal basis for **column space** of A .

The rest $n-k$ columns are null space of A^T (so that they form an orthogonal basis).

k

V is an $n \times n$ orthogonal matrix whose columns are the eigenvectors of $A^T A$.

The columns of V forms an orthogonal basis for **the whole input space**.

The **last $n-k$ columns** of V forms an orthogonal basis for **nullspace** of A .

The first k columns of V forms an orthogonal basis for row space of A .

D is a $n \times n$ diagonal matrix, with the non-negative eigenvalues of $A^T A$.

D' is a $m \times m$ diagonal, with the non-negative eigenvalues of $A A^T$.

D and D' have the same non-zero diagonal entries (except that the order might be different)

Conclusion

Any $m \times n$ matrix A can be factored as $A = U\Sigma V^T$, where:

$$\text{Rank}(A) = \text{Rank}(\Sigma) = k$$

Σ is a diagonal $m \times n$ matrix of the form:

$$\Sigma = \begin{bmatrix} \sigma_1 & & & & \\ & \ddots & & & \\ & & \sigma_k & & \\ & & & \ddots & \\ & & & & 0 \end{bmatrix}$$

$$\text{with } \sigma_1 \geq \sigma_2 \geq \dots \geq \sigma_k > 0$$


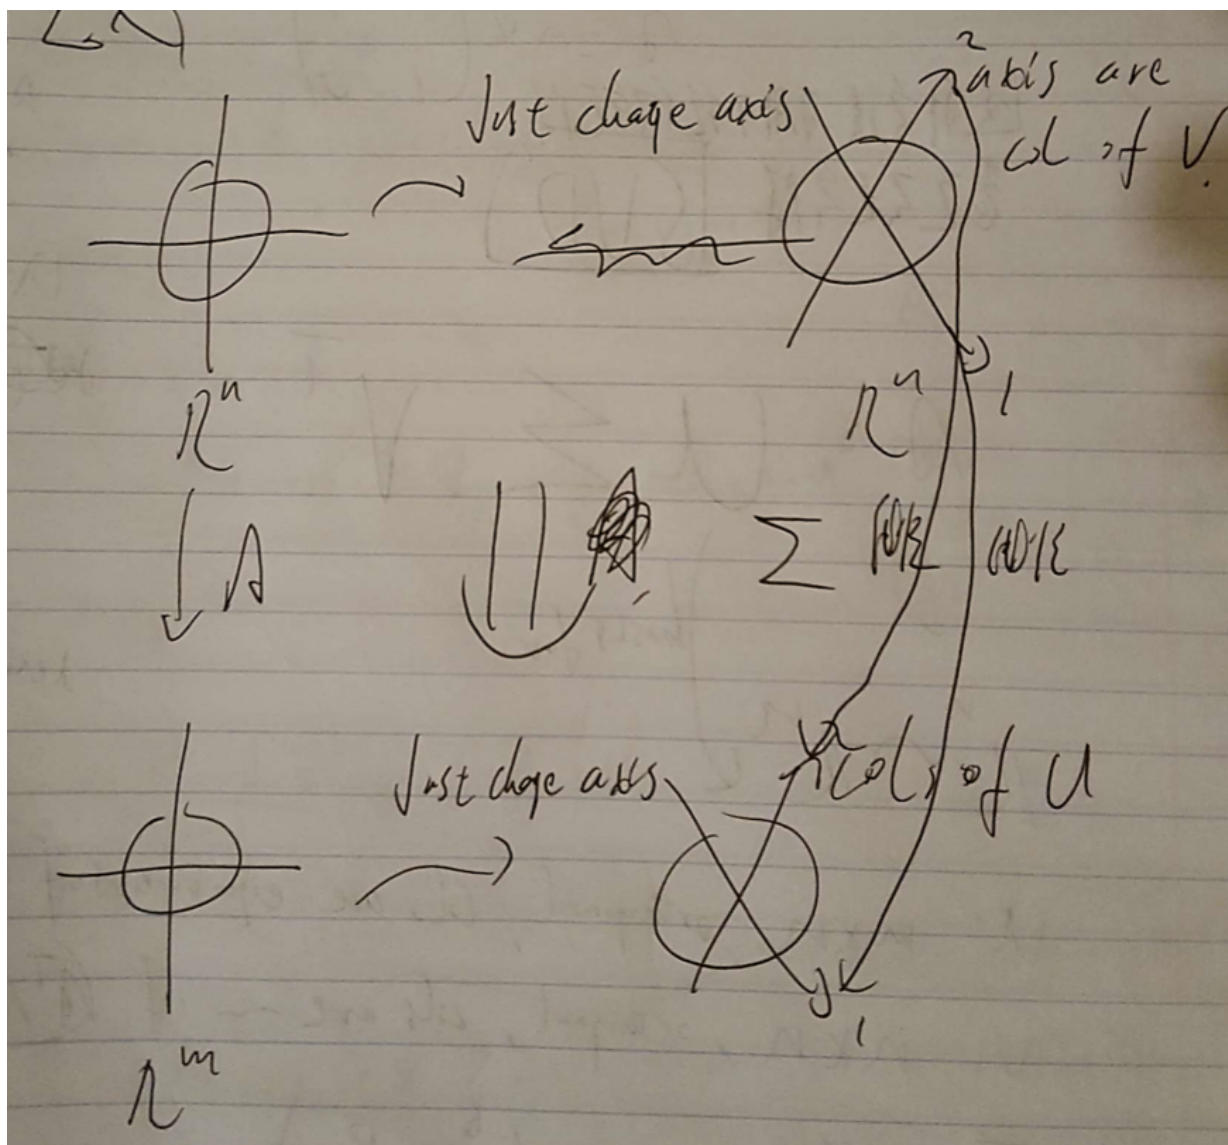
$$\text{and } k = \text{rank}(A)$$

$\sigma_i = \sqrt{\lambda_i}$, where λ_i is the nonzero eigenvalue of $A^T A$.

In fact, the nonzero eigenvalues of $A^T A$ are the same as of AA^T

When A is symmetric positive-definite, then SVD reduce to diagonalization.

Geometry




2.5.3 Solving linear systems

To solve $Ax=b$, we first calculate the SVD decomposition $A = U\Sigma V^T$, then compute $\bar{x} = V \frac{1}{\Sigma} U^T b$

what SVD done

The following sketch shows how SVD solve for $AX = b$ and $AX = b'$, where $b \in \text{colspace}(A)$, $b' \notin \text{colspace}(A)$



SVD solve $Ax=b$ by choosing x entirely from row space, which is closest to the origin among all possible solutions. $V \frac{1}{\Sigma} U^T$ is a pseudo-inverse.

SVD solve $Ax=b'$ by projecting b' onto the colspace of A , obtaining b then solve $Ax=b$. In other words, SVD obtain the least-square solution.

Pseudo-inverse: $Ax = b$, the pseudo-inverse solution to it is $x_p = (A^T A)^{-1} A^T b$.

$$(A^T A)^{-1} A^T = V \frac{1}{\Sigma} U^T$$

How SVD do it

Before we do it, we define $\frac{1}{\Sigma}$ as the diagonal matrix whose diagonal entries are of the form:

$$\left(\frac{1}{\Sigma} \right)_{ii} = \begin{cases} \frac{1}{\Sigma_{ii}} & \text{if } \Sigma_{ii} \neq 0 \\ 0 & \text{if } \Sigma_{ii} = 0 \end{cases}$$

Observe:

- if Σ is invertible, then $\frac{1}{\Sigma} = \Sigma^{-1}$
- if $\text{rank}(\Sigma) = k$ then

$$\frac{1}{\Sigma} \cdot \Sigma = \Sigma \cdot \frac{1}{\Sigma} = \begin{pmatrix} 1 & \boxed{} & 0 \\ \boxed{} & 1 & \boxed{} \\ 0 & \boxed{} & 0 \end{pmatrix}$$

Ok, now takes a look at SVD.

Firstly, $Ax = b \rightarrow$

$$U\Sigma V^T x = b$$

$$\Sigma V^T x = U^T b$$

$U^T b$ is the coordinate of b under an orthogonal basis, and the first k vectors of the basis forms the $\text{colspace}(A)$. In other word, $U^T b$ is the projection of b on $\text{Colspace}(A)$.

Thus, if $b \in \text{colspace}(A)$, $U^T b = [x_1, \dots, x_k, 0, \dots, 0]^T$

Next,

$$V^T x = \frac{1}{\Sigma} U^T b$$

$V^T x$ gives the coordinate of x under a basis whose last $n-k$ vectors spans $\text{Null}(A)$. In other words, x is projected on the row space of A ,

Considering the structure of $\frac{1}{\Sigma}$, it only do dilation on first k elements, then ignore the other elements.

Now, lets take an inverse look. what is the meaning of SVD solution \bar{x} ? What they do with b ?

clearly, it project b on $\text{colspace}(A)$, dilation, then project what remains onto $\text{rowspace}(A)$.

Now, lets take an look at what b could we get using a SVD solution \bar{x} .

$$\begin{aligned} A\bar{x} &= U\Sigma V^T V \frac{1}{\Sigma} U^T b \\ &= U \begin{bmatrix} 1 & & \\ & \ddots & \\ & & 0 \end{bmatrix} U^T b \end{aligned}$$

If $b \in \text{colspace}(A)$, then the matrix

$$\begin{bmatrix} 1 & & \\ & \ddots & \\ & & 0 \end{bmatrix}$$

actually do nothing, and $A\bar{x} = b$.

If $\mathbf{b} \notin \text{colspace}(\mathbf{A})$, then we only got \mathbf{b} projected on colspace, eliminated other elements, then project back to whole space. That is the projection of \mathbf{b} onto $\text{colspace}(\mathbf{A})$.